TANA: Fined-grained Side-channel Inforamtion Leakage Quantifications in Binaries

Anonymous

Abstract—Side-channel attacks allow attackers to infer some sensitive information based on non-functional characteristics. Existing works on address-base side-channel detection can provide a list of potential side-channels leakage sites. We observed that those works still have the following limitations: 1) Many software may have multiple information leakage sites. Some vulnerabilities could be more severe than others. But existing work couldnt tell the difference between those leakages. 2) An attacker could exploit multiple leakages at one time. However, no existing tool can report how much information is leaked in total.

To overcome the above limitations: we proposed a tool called TANA, which can not only find the side channels but can estimate how many bits are actually leaked through the leakage. TANA works in three phases. First, the application is executed to record the trace. Second, TANA runs the instruction level symbolic execution on the top of the execution trace. Tana will find side-channel information leakages and model each leakage as one unique math constraints. Finally, TANA will classify those constraints into independent multiple groups and run the multiple step monte carlo to estimate the information leakage. TANA will a very fined-grained result compared to existing tools.

We apply the tool on OpenSSL, MbedTLS and libjpeg and find several serious side channel vulnerabilities. We also evaluate the vulnerabilities from previous research. The result indicates most of the reported vulnerabilities are actually hard to exploit in practice.

I. INTRODUCTION

Side channels are inevitable in computer systems. The attacker can infer the sensitive information by observing execution behaviours. The software-based attackers(e.g., Cache-based side channels, controlled side attacks) are especially common because these attacks typically dont need any physical access. We observe that the root cause is that the implementation of the program may have different memory access patterns during the execution. Attackers can observe the pattern and infer the sensitive information based on source code.

Various countermeasures have been proposed to defend the address-based side-channel attacks. The basic idea is to identify and eliminate secret-dependent memory access patterns. It is very tedious and hard to manually find every the leakage site. To address the problem, a lot of methods have been proposed to automatically detect information leakage. While those tools can typically report a list of potential leakages sites, they still face a couple of limitations.

First, the existing tools [1] [2] can discover a list of potential leakage sites. But they fail to report how severe each potential leakage site could be. Most of leakages found by those tools are typically hard to exploit or leak very little information (e.g., one bit of the length of the key). It is worthwhile to

have a tool to tell how much information is actually leaks. For example, many secret dependent memory access patterns in OpenSSL have been known and studied for years. But the authors still doesnt fix them all because they think they arent not severe. The tool based on static analysis with abstract interpretation, can give leakage leakage upper bound, which is useful to justify the implementation is secure enough. But they cant indicate the leakage site is severe enough due to the over approximation. The dynamic method will take a concrete input and run the actual program. Those tools are typically very precise in term of true leakage. But none of those tools can actually report how many bits are actually leaked. For example, DATA [2] reports more than 2000 potential leakage sites for the RSA implementation of OpenSSL. But most of them were dismissed by the author after some manual inspections.

Second, many open source libraries may have multiple leakage sites. A side-channel attacker [3] can exploit multiple leakage sites at one time. The attacker may retrieve some information from one site and some other information from another site. It is hard to model how much information is actually leaked in total. Adding those leakages simply can get only an upper bound estimate of the total information leakage if those leakages arent independent. No existing tools can practically estimate the total information leakage from multiple leakage sites in open source libraries.

In the paper, we present a tool that can automatically identify and quantify address-based sensitive information leakage sites in real-world applications. The intuition is that most adversaries will exploit the different control-flow transfer and data-access patterns when the program process different sensitive data. We refer them as the potential information leakage sites. Our tool can provide a list of information leakage sites as well as how many bits they can leak if the attacker can observe one site, two sites or multiple sites. We define the amount of leaked information by the success rate of guessing the sensitive value after an adversary could increase after launching the side-channel attacks. It is interesting to mention that the definition here is different than the definition in several static analysis tools. We will explain the reason in the following sections. After that, we run the symbolic execution on the execution trace. We model each side-channel leakage as one logic formula. The sensitive input is divided into several independent bytes and each byte is regarded as the one symbol. Those formulas can unique model the side-channel vulnerability. In other words, if the application has a different sensitive input but still satisfy the formula, the code will still

leak the same information about the input. Those information leakage sites may spread in the whole program and their leakages arent dependent. Simply adding them up can only give the upper bound estimate. In order to accurately calculate the total information leakage, we must know the dependent relationships among those multiple leakages sites. Therefore, we introduce a monte carlo sampling method to estimate the total information leakage from the formula generated by previous symbolic execution.

We implement the proposed method as a tool named TANA that can precisely discover and quantify the information leakage vulnerabilities. We apply TANA on several crypto and non-crypto libraries including OpenSSL, MbedTLS and libjpeg. The experiments confirm that TANA can precisely identify the pre-known vulnerabilities, report how much information is leaked and which byte in the original sensitive buffer is leaked. The result shows that while those crypto libraries have a number of side-channels, most of them actually leak very little inforamtion. Also, we also use the tool to analysis the two reported side-channel attack in the libjpeg library. Finally, we present new vulnerabilities. With the help of TANA, we think those vulnerabilities are easily to be exploited.

In summary, we make the following contributions:

- We propose a novel trace-based analysis method that can quantify fined-grained side-channel information leakages. We model each information leakage vulnerability as one formula. Therefore, mutiple side-channel vulnerabilities can be seen as the disjunctive normal form (DNF) formula, which precisely models the program semantics.
- We transfer the information quantification problem into a probabilty distribution problem and use guided Monte Carlo sampling method to estimate the information leakage. Some initial results indicate the the sampling method suffers from the curse of dimensionality problem. We therefore design a guided sampling method and porvide the corresponding error esitimate.
- We implement the proposed method into a practical tool and apply it to a detect several software. TANA successfully identify the address-based side-channel vulnerabilities and provide the corresponding information leakge. The information lekage result provide the detailed information that help developers to fix the vulnerability.

II. BACKGROUND

In this section, we first present a basic introduction about the memory-based side-channel attack. Those attacks are exactly what we attempt to study in the paper. After that we will present existing work on side-channel detection and quantification. We will also analyze strengths and limitations of those quantification methods.

A. Address-based Side-Channels

Address-based side-channels are information channels that can leak sensitive information unintended through the different behaviors when the program accesses different memory addrsss. Fundamentally, those difrences were caused by the memory hierarchy design in modern computer systems. When the CPU fetches the data, it will first search the cache, which stores copies of the data from the frequently used main memory. If the data doesn't exist in the cache, the CPU will read the data from the main memory (RAM). Classified by the layer caused the side-channel, we introduce two kinds of commom side-channels: cache-based side-channel attacks and memory-based side-channel attacks.

1) Cache-based Attack: In general, the cached-based side-channel attacks seek information rely on the time differences between the cache miss and cache hit. Here we introduce two types of cache attacks: PRIME+PROBE, FLUSH+RELOAD.

PRIME+PROBE targets a single cache set. It has two phases. During the "prime" phase, the attacker fills the cache set will his own data. In the second "probe" phase, the attacker accesses the cache set again. If the victim accesses the cache set and evicts part of the data, the attacker will experience a slow mesurement. If not, the mesurement will be fast.

FLUSH+RELOAD targets a single cache line. It requires the attacker and victim share the same memory. It also have two phases. During the "flush" phase, the attacker will flush the "monitered memory" from the cache. Then the attacker wait for the victim to access the memory. In the third phase, the attacker reload the "monitered memory". If the time is short, which indicates there is a cache hit and the victim reolads the memory before. On the other hand, the time will be longer since the CPU need to reolad the memory into the cache line.

2) Memory-based Attack: Memory-based side-channel attack [] exploits the different behaviors when the program accesses different page tables. The controlled-channel attack [4], which works in the kernel space, can infer the sensitive data in the shielding systems by observing the page fault sequences by restricting some code and data pages.

After examing the memory-based side-channels attack. We find the fundamentally reason of those attacks are due to secret-dependent memory access and control flow transfers.

```
void sample_encrypt_setup                              1
            (encrypt_ctx *ctx,                         2
             const unsigned char *key)                 3
{                                                      4
    int i, j, k;                                       5
    unsigned char *m, T[256];                          6
    m = ctx->m;                                        7
    for(i = 0 ; i < 256; i++, k++)                     8
    {                                                  9
        // Secret Dependent Memory Access             10
        m[i] = Table[key[k] % 256];                   11
        // Secret Dependent Control Flow Transfer     12
        if(key[k] == 0)                               13
        {                                             14
            m[i] = k * key[i];                        15
            m[i] = m[i] % 256;                        16
        }                                             17
    }                                                 18
                                                      19
}                                                     20
```

Listing 1. Sample code shows secret-dependent memory access and secret-dependent control-flow transfer.

For exampls, the above code 1 show an simple encryption function that has the two kinds of side-channels. At the line

11, dependending on the value of key, the code will access the different entry in the predefined table /textbfTable. At the line 13, the code will do a series of computation and determine if the code in the if branch is executed or not. Such vulnerabilities could leak to the memory-based side-channles. We identify and quantify the leakage of the two kinds of vulnerabilities in the paper.

B. Information Leakage Quantification

Given an event e which occurs with the probability $P(e)$, if the event e happens, then we receive

$$I = -log(P(e)) \tag{1}$$

bits of information by knowing the event e.

The above definition is obvious. Suppose a char variable a in C program has the size of one byte (8 bits), so the value in the variable can range from 0 - 255. We assume the a has the uniform distribution. If at one time we observe the a equals to 1, the probability will be 1/256. So the information we get is $-log(1/256) = 8 bits$, which is exactly the size of the char variable in C program.

Existing works on information leakage quantification are based on mutual information or min-entropy [5]. In their frameworks, the input sensitive information K is viewed as random variables. Let k_i be one of the possible value of K. The Shannon entropy $H(K)$ is defined by

$$H(K) = - \sum_{k_i \in K} P(k_i)log(P(k_i)) \tag{2}$$

The Shannon entropy can be used to quantify the initial uncertainty about the sensitive information. Suppose a program (P) with the K as the sensitive input, an adversary has some observations (O) through the side-channels. In this work, the observations are referred to the secret-dependent control-flows and secret-dependent data-accesses patterns. The conditional entropy $H(K|O)$ is

$$H(K|O) = - \sum_{o_j \in O} P(o_j) \sum_{k_i \in K} P(k_i|o_j)log(P(k_i|o_j)) \tag{3}$$

Intuitively, the conditional information marks the uncertainty about K after the adversary has gained some observations (O).

Many previous works use the mutual information $I(K;O)$ to quantify the leakage which is defined as follows:

$$Leakage = I(K;O) = \sum_{k_i \in K} \sum_{o_j \in O} P(k_i, o_j)log(\frac{P(k_i, o_j)}{P(k_i)P(o_j)}) \tag{4}$$

where $P(k_i, o_i)$ is the joint discrete distribution of K and O. Alternatively, the mutual information can also be computed with the following equation:

$$Leakage = I(K;O) = H(K) - H(K|O) = H(O) - H(O|K) \tag{5}$$

For a deterministic program, once the input K is fixed, the program will have the same control-flow transfers and data-access patterns. As a result, $P(k_i, o_j)$ will always equals to 1. So the conditional entropty $H(O|K)$ will equal to zero.

Now the leakage defined by the mutual information can be simplified into:

$$Leakage = I(K;O) = H(O) \tag{6}$$

In other words, once we know the distribution of those memory-access patterns. We can calculate how much information is actually leaked.

Another common method is based on maximal leakage [5]–[7].

$$Leakage = log(C(O)) \tag{7}$$

$C(O)$ represents the number of different observations that an attacker can have.

Now we provide a concrete example to show how the two types of quantification definition works.

```
unsigned char key = input();              1
// key = [0 ... 255]                       2
if(key = 128)                              3
    A(); // branch 1                       4
else if (key < 64)                         5
    B(); // branch 2                       6
else if (key < 128)                        7
    C(); // branch 3                       8
else                                       9
    D(); // branch 4                      10
```

Listing 2. A simple program

Maximal leakage Dependenting on the value of key, the code can run four different branches which corrosponding to four different observations. Therefore, by the maximal leakage definition, the leakage equals to $log 4 = 2$ bits.

Mutual Information If the key satisfies the uniform distribution, the probability of the code runs each branch can be computed with the following result: Therefore, the leakage

Branch	A	B	C	D
P	1/256	64/256	64/256	127/256

TABLE I
THE DISTRIBUTION OF OBSERVATIONS

equals to $\frac{1}{256}log\frac{1}{256} + \frac{1}{4}log\frac{1}{4}*2 + \frac{127}{256}log\frac{127}{256} = 1.7$ bits.

III. THREAT MODEL

We consider an attacker who shares the same hardware resource with the victim. The attacker attempt to retrieve sensitive information via memory-based side-channel attack. The attacker has no direct access to the memory or cache, but can probe the memory or cache at each program point. Similar to DATA, the attacker will face many noisy observations or can only observe a limited of memory or caches in practice. For the project, we assume the attacker can have noise-free observations. This threat model captures most of the cache-based and memory-based side channel attacks. We only consider the deterministic program for the project.

IV. MOTIVATION

In this section, we articulate several chanllenges and existing problems in quantifying the side-channel vulnerability leakages. We briefly describe the chanllenges and then present the corresponding solutions.

A. Information Leakage Definition

Existing static-based side-channel quantification works defined information leakage as the mutual information or the max leakage. Those definitions provide strong security guarantee when trying to show a program is secure if the their methods say the program leaks zero bits of information.

However, the above definition is less useful when if the program has some leakages. Considering the example in section 2, if an attacker observes the code runs branch A, the attacker can know the key actually equals to 128. Suppose it is a dummy password checker, in which case the attacker can fully retrieve the password. Therefore, the total information leakage should be 8 bits, which equals to the size of unsigned char. According to the mutual definition, however, the leakage will be 1.7 bits. The maximal information leakage is 2 bits. Both approaches fail to precisely tell how much information is actually leaked during the exectution.

The problem with the existing method is that they are static-based and the sensitive input values are neglected. They assume the attacker runs the program multiple times with many different sensitive information as the input. Both the mutual information and the max-leakage give an "average" estimate of the information leakage. But it isn't the typical scenario for an adversary to launch the side-channel attack. When a side-channel attack happens, the adversary wants to retrieve the sensitive information. So it is very likely that the sensitive input is fixed (e.g. AES keys). The adversary will run the attack over and over again and guesses bit by bit. Like the previous example, the existing static method doesnt work well in those situations.

In the project, we hope to give a very precise definition of information leakages. Suppose an attacker run the target program mutiple times with one fixed input, we want to how much information he can infer by oberserving the memory access patterns. We come to the simple slogan [5] where the information leakage equals:

initial uncertainty - remaining uncertainty

If an adversary has zero knowledge about the input before the attack. The initial uncertainty equals to the size the input. As for the remaining uncertainty, we come to the original definition of the information content.

Solutioin: We quantify the information leakage with the following definition.

Definition 1. *Given a program P with the input set K, an adversary has the observation o when the input $k \in K$. We denote it as*

$$P(k) = o$$

The leakage L_{Pko} based on the oberservation is

$$L_{Pko} = log_2|K| - log_2|K^o|$$

where

$$K^o = \{k' | k' \in K \ and \ P(k') = o\}$$

With the new definition, if the attacker observes the branch 1 was executed, then the $K^{o^1} = \{128\}$. Therefore, the

information leakage $L_{Pko^1} = log_2 256 - log_2 1 = 8$ bits, which means the key is totally leaked. If the attacker observes the code runs other branches, the leaked information is shown in the following table.

Branch	1	2	3	4
K^o	1	64	64	127
L_{Pko}(bits)	8	2	2	1

TABLE II
THE LEAKED INFORMATION BY THE DEFINITION

B. Multiple Leak Sites

Real-world software can have various side-channel vulnerabilities. Those vulnerabilities may spread in the whole program. An adversary may exploit more than one side-channel vulnerabilities to gain more information [3], [4]. For example, the controlled-side channel attack [4], the author demonstrates an attack against a popular spell checking tool, Hunspell. By observing four sets of secret-dependent memory accesses sites in two functions $HashMgr :: addword$ and $HashMgr :: lookup$, the author can recover the word that Hunspell checks.

For the Hunspell, the attacker manually studies the source code of Hunspell, figure out the relation of those vulnerabilities and launch the attack. In order to precisely quantify the total information leakage, we need to know the relation of those leakage sites.

```
1  unsigned char k1, k2, k3;
2  ...
3  t1 = T[k1 % 8];          // Leakage 1
4  t2 = T[k2 % 8];          // Leakage 2
5  t3 = T[k3 % 8];          // Leakage 3
6  bool flag = foo(t1, t2, t3);
7  if(flag) {
8      if(k1 > 128)         // Leakage 4
9          A();
10 }
11 else {
12     if(k2 > 128)         // Leakage 5
13         B();
14 }
15 if(k1 + k2 + k3 > 128) // Leakage 6
16     C();
```

Listing 3. Multiple leakages

Considering the running example in 3, in which $k1$, $k2$ and $k3$ are the sensitive key. The code has six different leakage. Leakage 1, 2, 3 are the secret-dependent data accesses and leakage 4, 5, 6 are the secret-dependent control-flow transfers. However, it is very hard to estimate total leakages. For example, the attacker can infer the last three digits of $k1$, $k2$, $k3$ from leakage 1, 2, 3. So those leakages are independent. For leakage 1, 4, 6, however, we have no idea about the total information leakage.

For a real program, it is tough to estimate the total information leakage for the following reasons. First, the real-world applications have more than thousands of lines of code. Side-channel vulnerabilities could exist in many different functions of the source code. One leakage site leaks the temporary value. But the value contains some information about the original

buffer. It is hard to know how the the sensitive value affects the temporary value. Second, some leakages sites may be dependent. The occurrence of the first affects the occurrence of the second sites. We can simply add them up. Third, leakage sites are in the different block of the control-flow graph, which means that only one of the two leakages site may execute during the execution.

Suppose one program has two side-channel vulnerabilities A and B, which leaks L_A and L_B bits during the execution. The total leaked information is noted as L_{Total}. The relation between A and B have the following three cases.

1) Independent Leakages: If A and B are independent leakages, the total information leakage will be:

$$L_{total} = L_A + L_B \tag{8}$$

2) Dependent Leakages: If A and B are dependent leakages, the total information leakage will be:

$$\max\{L_A, L_B\} <= L_{total} < L_A + L_B \tag{9}$$

3) Mutual Exclusive Leakages: If A and B are mutual exclusive leakges, then only A or B can be observed for one fixed input. The total information leakage will be L_A or L_B.

According to above definition, leakage 1, 2, 3 are independent leakages. Leakage 4, 5 are mutual exclusive leakages.

We run the symbolic execution on the top of the execution traces. At the beginning of the exectution, each byte in the sensitive buffer is modeled with with a symbol. After that, the symbolic execution engine interprets each instruction of the execution traces. So every values in the registers or memory cells is modeled with a math formula.

Given a program P, k is the sensitive input. The k should be a value in a memory cell or a sequential buffer (e.g., an array). We use k_i to denote the sensitive information, where i is the index of the byte in the original buffer. We can have the following equations. The t1, t2, t3, is the temporary values during the execution.

$$t_1 = f_1(k_1, k2, ...k_n)$$

$$t_2 = f_2(k1, k2, ...k_n)$$

$$t_3 = f_3(k1, k2, ...k_n)$$

...

$$t_m = f_m(k1, k2, ...k_n)$$

After that, we model each potential leakage sites as a math formulas.

The attacker can retrieve the sensitive information by observing the different patterns in control-flows and data access when the program process different sensitive information. We refer them as the secret-dependent control flow and secret-dependent data access accordingly.

4) Secret-dependent Control Flow: Here is an example of the secret-dependent control-flows. Consider the code snippet in List 1. Here the key is the confidential data. The code will have different behaviours (time, cache access) dependenting on which branch is actually executing. By observing the behaviour, the attacker can infer which branch actually executed and know some of the sensitive information. One of the famous leakage example is the square and multiply in many RSA implementations.

For example, the attacker knows the key equals to zero if he observes the code run the branch1. Because key has 256 different possibilities. The original key has lg256 = 8 bits information. If the attacker can observe the code run branch 1. Then he will knows the key equals to zero. If the code run branch 2, the attacker can infer the key doesnt equal to zero.

```
Branch 1
temp = 0xb;
0 =< key <= 256;
temp = key/2;
```

Information Leakage = -log(1/p) = -log(1/256) = 8 bits

```
Branch 2
temp != 0;
0 =< key <= 256;
temp = key/2;
```

Information Leakage = -log(255/256) bits

```
T[64]; // Lookup tables with 64 entries
index = key % 63;
temp = T[index];
// Secret-dependent memory access
```

The simple program above is an example of secret dependent memory access. Here T is a precomputed tables with sixty-four entries. Depending on the values of key, the program may access any values in the array. Those kind of code patterns may wildly exist in many crypto and media libraries.

Suppose the attackers observe the code accesses the first entry of the lookup tables. We can have the following formulas.

```
key mod 63 = 1
0 =< key <= 256
```

So the key can be one of the following values: 1 64 127 190 253
Information leakages = -log(5/256) = 5.6 bits

C. Scalability and Performance

After we transfer each potential leaks sites into logic formula. We can group several formulas together to estimate the total information leakage. One naive way is to use the Monte Carlo sampling estimate the number of input keys. With the definition 1, we can estimate the total information leakage.

However, some pre-experiments show that above approach suffers from the unberable cost, which impede its usage

to detect and quantify side-channel leakages in real-world applications. We systematically analyze the performance bottlenecks of the whole process. In general, the performance suffers from the two following reasons.

- Symbolic Execution.
- The Naive Monte Carlo Sampling.

1) Symbolic Execution: Symbolic execution interprets each instructions and update the memory cells and registers with a formula that captured the sementics. Unfortunately, the number of machine instructions are huge and the sementics of each instruction is complex. For example, the Intel devloper manual [8] introduces more than 1000 different X86 instructions. It is tedious to manually implement the rules for every instructions.

Therefore, existing binary analysis tools [9], [10] will translate machine instructions into intermediate languages (IR). The IR typically has fewer instructions compared to the original machine instructions. The IR layer designs, which significantly simplify the implementations, also introduce significant overhead as well [11].

2) Monte Carlo Sampling: For an application with n bits secret, there are total 2^n possible inputs. Of the 2^n possible secrets,

V. DESIGN

In this section, we will explain the design decisions to realize A.

A. Trace Logging

The trace information can be logged via some emulators (e.g., QEMU) or Binary Instrumentation Tools. For our project, we write an Intel Pin Tool to record the execution traces. The trace data has the following information: Each instruction mnemonics and its memory address The operands of each instruction and their values The memory address of sensitive information and its length The value of eflags register Most software developers stores sensitive information in an array, a variable or a buffer, which means that those data is stored in a contiguous area in the memory. We use the symbol information in the binary to track the address in the memory.

B. Instruction Level Symbolic Execution

The main purpose of the this step is to generate constraints of the input sensitive information for the execution trace. If we give the target program a new input which is different from the origin input that was used to generate the execution trace but still satisfies those constraints, the new execution trace will still have the same control flow and data access patterns.

The tool runs the symbolic execution on the top of the execution traces. At the beginning of the symbolic execution, the tool creates fresh symbols for each byte in the sensitive buffer. For other data in the register or memory at the beginning of the symbolic execution, we use the concrete value from the runtime information collected in the previous step. During the symbolic execution, the tool will maintain a symbol and a concrete value for every variables in the memory and registers. The formula is made up with concrete values and the input key as the symbols calculated through the symbolic execution. For each formula, the tool will check weather it can be reduced into a concrete values. If so, the tool will only use the concrete values in the following symbolic execution.

C. Verification and Optimization

We run the symbolic execution on the top of x86 instructions to achieve the better performance and accuracy of the memory model. In other words, we dont rely on any intermediate languages to simplify the symbolic execution. While the implementation itself has a lot of benefits, we need to implement the symbolic execution rules for more than one thousand x86 instructions. However, due to the complexity of X86, it is inevitable to make mistakes. Therefore, we verify the correctness of the symbolic execution. The tool will collect the runtime information (Register values, memory values) and compare them with the values generated from the symbolic execution. Whenever the tool finishes the symbolic execution of each instruction, the tool will compare the formula for each symbol and its actual value. If the two values dont match, we check the code and fix the error. Also, if the formula doesnt contain the any symbols, the tool will use the concrete value instead of symbolic execution.

D. Secret-dependent control-flows

An adversary can infer sensitive information from secret dependent control-flows. There are two kinds of control-transfer instructions: the unconditional control transfer instructions and the conditional transfer instructions. The unconditional instructions, like CALL, JUMP, RET transfer control from one code segment location to another. Since the transfer is independent from the input sensitive information, an attacker was not able to infer any sensitive information from the control-flow. So the unconditional control transfer doesnt leak any information based on our threat model. During the symbolic execution, we just update the register information and memory cells with the new formulas.

The conditional transfer instructions, like conditional jump, may or may not transfer control, depending the CPU state. For any conditional jump, the CPU will test if certain condition flag (e.g., CF = 0, ZF =1) is true or false and jump to the certain branches respectively. So the symbolic engine will compute the flag and represent the flag in a symbol formula. Because we are running on a symbolic execution on a execution trace, we know which branch executes. If a conditional jump uses the CPU status flag, we will generate the constraint accordingly.

For examples,

```
...
0x0000e781          add dword [local_14h], 1
0x0000e785          cmp dword [local_14h], 4
0x0000e789          jne 0xe7df
0x0000e78b          mov dword [local_14h], 0
```

At the beginning of the instruction segment, the value at the address of local14h can be written as F(K). At the address e785, the value will be updated with F(K)+1. Then the code compare the value with 4 and use the result as a conditional jump. Based on the result, we can have the following constrain:

F(K) + 1 = 4

E. Secret-dependent data access

Like control-flows, an adversary can also infer sensitive information from the data access pattern as well. We try to find this kind of leakages by checking every memory operand of the instruction. We generate the memory addressing formula. As discussed before, every symbols in the formula is the input key. If the formula doesnt contain any symbols, the memory access is independent from the input sensitive information and wont leak any sensitive information according to our threat model. Otherwise, we will generate the constraint for the memory addressing.

Monte Carlo Volume Sampling From the above steps, we already have the constraints from the execution trace. The only variables in those constraints is the sensitive data. An adversary who wants to infer the sensitive data based on side-channel attacks cant observe the sensitive information directly. The adversary, however, can observe the memory access pattern of the software.

We use the Monte Carlo Sampling to calculate the ratio of input that satisfies those constraints. We will estimate how much information is actually leaked from each input key.

Normalizations The goal of the normalizations is to simplify the constraints. Each constraint will be evaluated multiple times during the following sampling, we would like to make those formula simpler to reduce the whole execution time. Each formula is implemented as a abstract syntax tree. We apply a series normalization rules (e.g. key1 xor key1 = 0) to simplify the formula.

Key1 + 1 + 2 + key3 ¡ key2 =¿ key1 + 3 ¡ key2

Split the independent constraints into multiple groups Each constraints may have multiple symbols as the input. If each two formulas have complete different inputs, then those leakages modeled by the constraints are independent.

Multiple Stage Monte Carlo Sampling After the above steps, each side channel leakage vulnerability is transfered into a formula. The only symbol in the orginal formula is the key. An adversary who wants to infer the sensitive data based on side-channel attacks cant observe the sensitive information directly.

The adversary, however, can observe the memory access pattern of the software. If the attacker can observe one leakage site, the attack is modeled as one formula. If he observe mutiple information leakage, the attack is modeled as DNF formula. Calculating the total information leakage can be reduced to the problem of approximating the number of solutions. One intution way is to use the Monte Carlo method. However, the number of satisfying keys could be exponentially small. Image an attacker who can recover the one unique key after the attacks, in such case, the DNF formula F = (f(k)) will only have one solution.

In order to estimate how much information is actually leaked from each input key. we have to calculate the ratio of input that satisfies those constraints. One intuition is to use the Monte Carlo method to approximate the number of solutions.

VI. DISCUSSION ON IMPLEMENTATION

A. Binary code vs source code

Many existing works find side-channels vulnerabilities from source code level or intermediate languages (e.g., LLVM IR). Those approaches will have the following questions. First, many compilers can translate the operator into branches. For example, the GCC compiler will translate the ! operator into conditional branches. If the branch is secret-dependent, the attacker could learn some sensitive information. But those source-based methods will fail to detect those vulnerabilities. Second, some if else in the source code can be converted into single conditional instructions (e.g., cmov). Source-based methods will still regard them as potential leakages, which will introduce false positives.

B. X86 vs Intermediate Languages

The tool takes in an unmodified binary and the marked sensitive information as the input. The tool will first start with logging the execution trace. Then the tool will symbolizes the sensitive information into multiple symbols and start with the symbolic execution. During the symbolic execution, the tool will find any potentials leakage sites as well as the path constraints. For each leakage site, whether it is the tool will also generate the constraint. Then, the tool will split the constraints into multiple group. After that, the tool will run monte carlo sampling to estimate the total information leakages.

REFERENCES

[1] S. Wang, P. Wang, X. Liu, D. Zhang, and D. Wu, "Cached: Identifying cache-based timing channels in production software," in *26th USENIX Security Symposium (USENIX Security 17)*. Vancouver, BC: USENIX Association, 2017, pp. 235–252. [Online]. Available: https://www.usenix.org/conference/usenixsecurity17/technical-sessions/presentation/wang-shuai

[2] S. Weiser, A. Zankl, R. Spreitzer, K. Miller, S. Mangard, and G. Sigl, "DATA – differential address trace analysis: Finding address-based side-channels in binaries," in *27th USENIX Security Symposium (USENIX Security 18)*. Baltimore, MD: USENIX Association, 2018, pp. 603–620. [Online]. Available: https://www.usenix.org/conference/usenixsecurity18/presentation/weiser

[3] D. Gruss, R. Spreitzer, and S. Mangard, "Cache template attacks: Automating attacks on inclusive last-level caches," in *24th USENIX Security Symposium (USENIX Security 15)*. Washington, D.C.: USENIX Association, 2015, pp. 897–912. [Online]. Available: https://www.usenix.org/conference/usenixsecurity15/technical-sessions/presentation/gruss

[4] Y. Xu, W. Cui, and M. Peinado, "Controlled-channel attacks: Deterministic side channels for untrusted operating systems," in *2015 IEEE Symposium on Security and Privacy*, May 2015, pp. 640–656.

[5] G. Smith, "On the foundations of quantitative information flow," in *Foundations of Software Science and Computational Structures*, L. de Alfaro, Ed. Berlin, Heidelberg: Springer Berlin Heidelberg, 2009, pp. 288–302.

[6] B. Kopf, L. Mauborgne, and M. Ochoa, "Automatic quantification of cache side-channels," in *Computer Aided Verification*, P. Madhusudan and S. A. Seshia, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 564–580.

[7] G. Doychev, D. Feld, B. Kopf, L. Mauborgne, and J. Reineke, "Cacheaudit: A tool for the static analysis of cache side channels," in *Presented as part of the 22nd USENIX Security Symposium (USENIX Security 13)*. Washington, D.C.: USENIX, 2013, pp. 431–446. [Online]. Available: https://www.usenix.org/conference/usenixsecurity13/technical-sessions/paper/doychev

[8] Intel Corporation, *Intel$^{\circledR}$ 64 and IA-32 Architectures Software Developer's Manual*, December 2009, no. 253669-033US.

[9] Y. Shoshitaishvili, R. Wang, C. Salls, N. Stephens, M. Polino, A. Dutcher, J. Grosen, S. Feng, C. Hauser, C. Kruegel, and G. Vigna, "SoK: (State of) The Art of War: Offensive Techniques in Binary Analysis," in *IEEE Symposium on Security and Privacy*, 2016.

[10] D. Brumley, I. Jager, T. Avgerinos, and E. J. Schwartz, "Bap: A binary analysis platform," in *Computer Aided Verification*, G. Gopalakrishnan and S. Qadeer, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 2011, pp. 463–469.

[11] I. Yun, S. Lee, M. Xu, Y. Jang, and T. Kim, "QSYM : A practical concolic execution engine tailored for hybrid fuzzing," in *27th USENIX Security Symposium (USENIX Security 18)*. Baltimore, MD: USENIX Association, 2018, pp. 745–761. [Online]. Available: https://www.usenix.org/conference/usenixsecurity18/presentation/yun