TBD: Precise and Practical Sampling-based Side channel Information Leakage Quantification

Bao Qinkun

What is a side-channel attack ?

Wikipedia: A side-channel attack is any attack that allows attackers to infer some sensitive information based on non-functional characteristics.

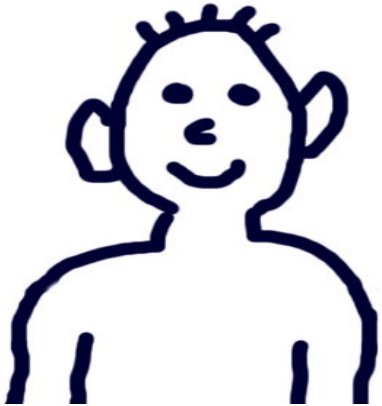

A naïve Ph.D. student

A prestigious professor

Side-channel attacks

Secret: whether your advisor is in the office

What you can do:

1. Ask your advisor (Not a good idea)

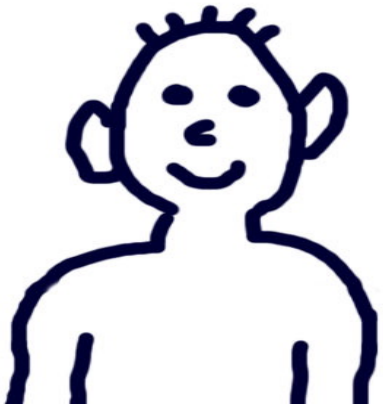

Attacker

Victim

Side-channel attacks

Secret: whether your advisor is in the office

What you can do:

1. Ask your advisor (Side effect)

Insight:

Your advisor always drives to work.

The parking lot is shared with you.

Side-channel attacks

Secret: whether your advisor is in the office

What you should do:

1. Go to the parking lot
2. Find your advisor's vehicle

Address-based Side Channels

Side-channels in computer systems: timing, power, CPU usage, cache …

```
unsigned char key = input();
// 4 bits information
unsigned char temp = key/2;
if (temp = 0xb)
// branch 0 execution 2s
else
// branch 1 execution 1s
```

Secret-dependent control-flow transfers (CF)

```
T[1024];
// Lookup tables with 1024 entries
index = key % 1024;
temp = T[index]
// Secret-dependent memory access
```

Secret-dependent memory accesses (DATA)

Cache-based: Prime + Probe, Flush + Reload, Flush + Evict
Memory-based: Controlled-channel attack, page-table attack

Side-channel Vulnerability Detections

- Software countermeasures

- Find vulnerabilities
  - It's hard. Real cryptosystems are complicated
  - Those vulnerabilities are architecture-dependent
  - Existing works: CacheAudit, CacheD, DATA, CacheS

- Fix vulnerabilities
  - It's also hard. Side-channel vulnerabilities are inevitable.
  - Many vulnerabilities are unpatched.

Motivation

	Algorithm	Differences	Generic			Specific	
			Dismissed	CF	Data	Byte/Bit	HW
OpenSSL	AES-NI	0 (2)	0	0	0 (2)	0 (2)	-
	AES-VP	0	0	0	0	0	-
	AES bit-sliced	4	0	0	4	4	-
	AES T-table	20	0	0	20	20	-
	Blowfish	194	0	0	194	171	-
	Camellia	82	0	0	82	55	-
	CAST	202	0	0	202	133	-
	DES	138	0	0	138	63	-
	Triple DES	410	0	0	410	292	-
	ECDSA (secp256k1)	515	487	1	27	3	1
	DSA	781	354	160	267	19	33
	RSA	2248	1510	278	460	11	139
PyCrypto	AES	96	0	0	96	96	-
	ARC4	5	0	0	5	5	-
	Blowfish	384	0	0	384	384	-
	CAST	284	0	0	284	216	-
	Triple DES	108	0	12	96	101	-

Table 1. Leakage summary of algorithms

S.Weiser,A.Zankl,R.Spreitzer,K.Miller,S.Mangard, andG.Sigl.
DATA - Differential Address Trace Analysis: Finding Address-based Side-Channels in Binaries. In *USENIX Security*, 2018.

Motivation

Existing works:

- False positives
  - e.g., 2248 potential leakages for RSA in OpenSSL, 1510 of them were dismissed
- Report many insensitive leakages
- Only on one leakage site

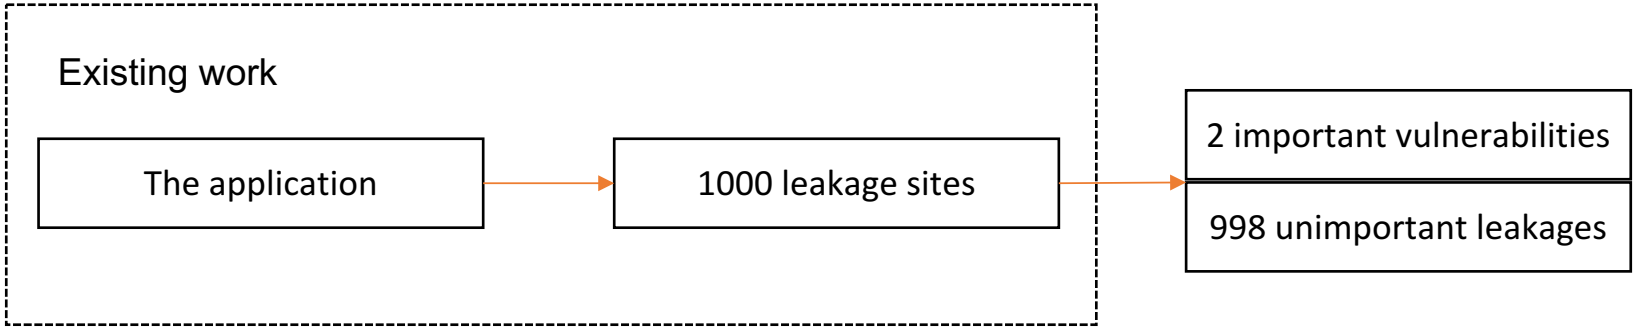

Our Objective: Identify and quantify address-based side-channels precisely.

Quantification

- Information Theory

The Bell System Technical Journal

Vol. XXVII *July, 1948* *No. 3*

A Mathematical Theory of Communication

By C. E. SHANNON

Introduction

THE recent development of various methods of modulation such as PCM and PPM which exchange bandwidth for signal-to-noise ratio has intensified the interest in a general theory of communication. A basis for such a theory is contained in the important papers of Nyquist[1] and Hartley[2] on this subject. In the present paper we will extend the theory to include a

Eight Apples

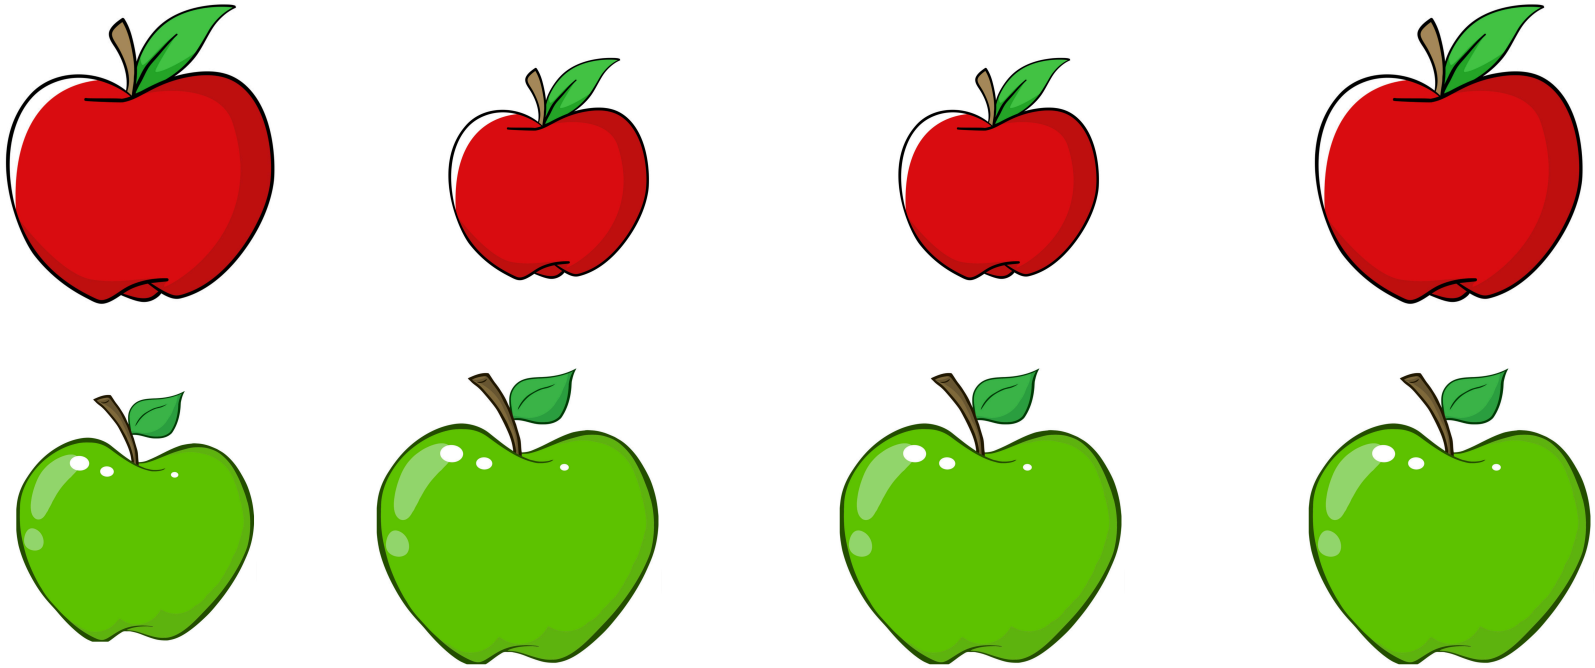

Secret: which apples is poisonous?

Eight Apples

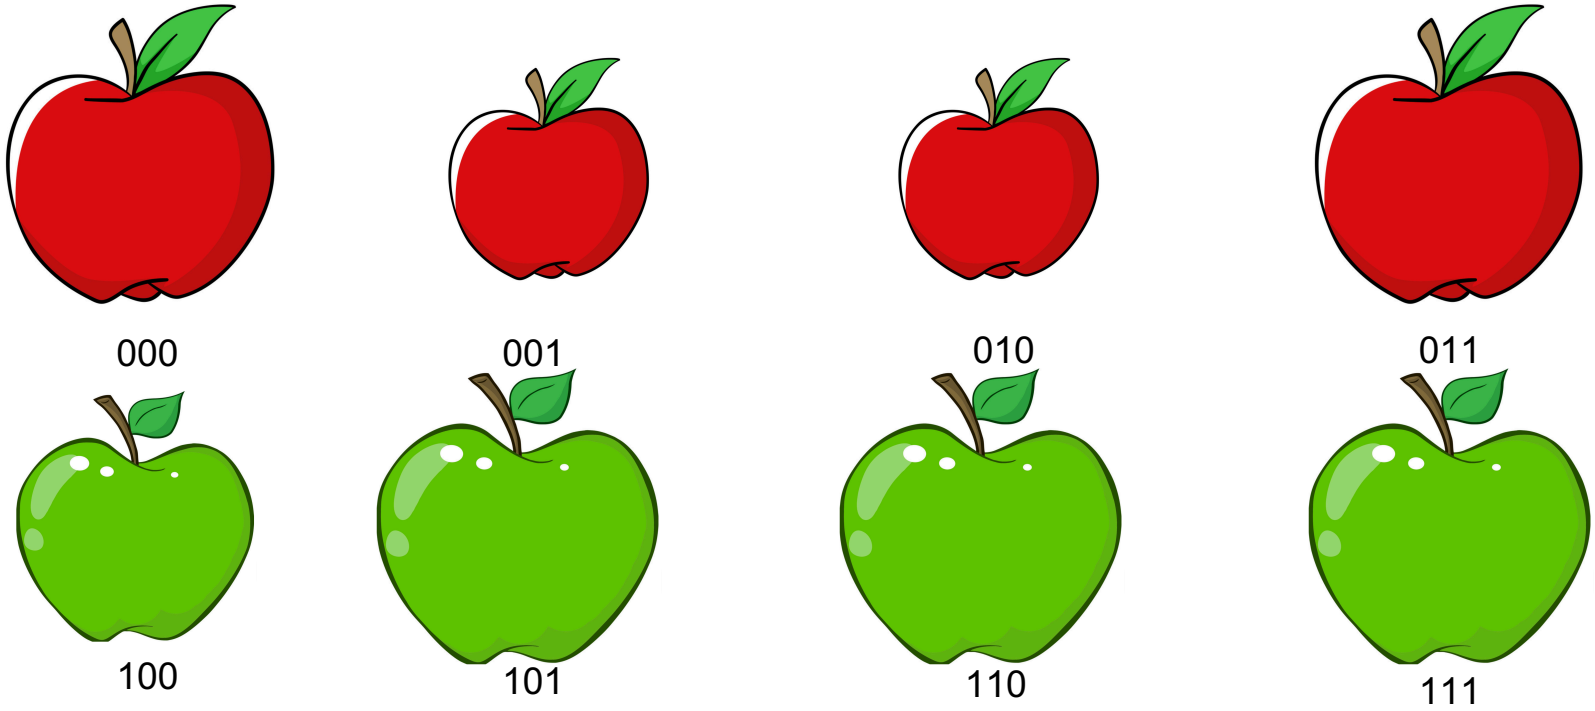

Secret: which apples is poisonous?

Eight Apples

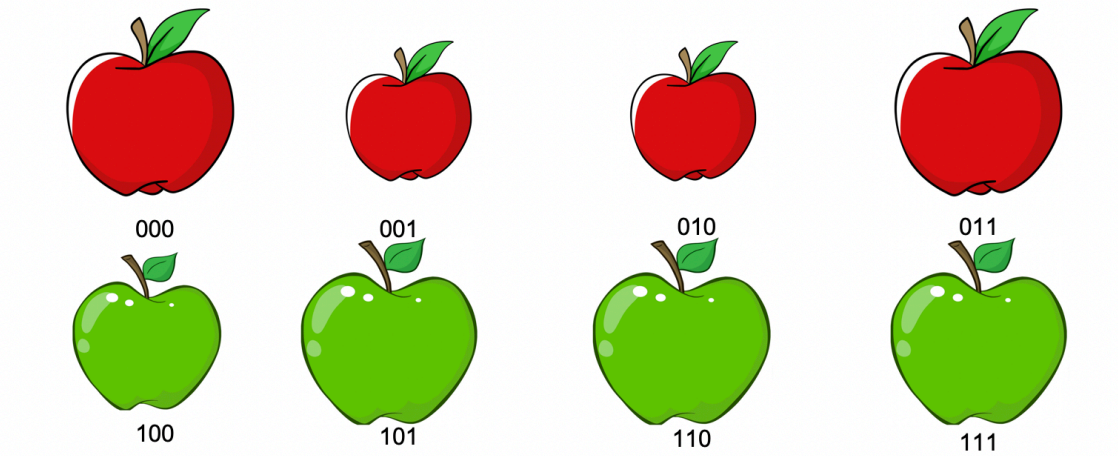

Secret:
Which apple is poisonous?

Before:
Each apple can be poisonous.
So we have 3 bits information.

Eight Apples

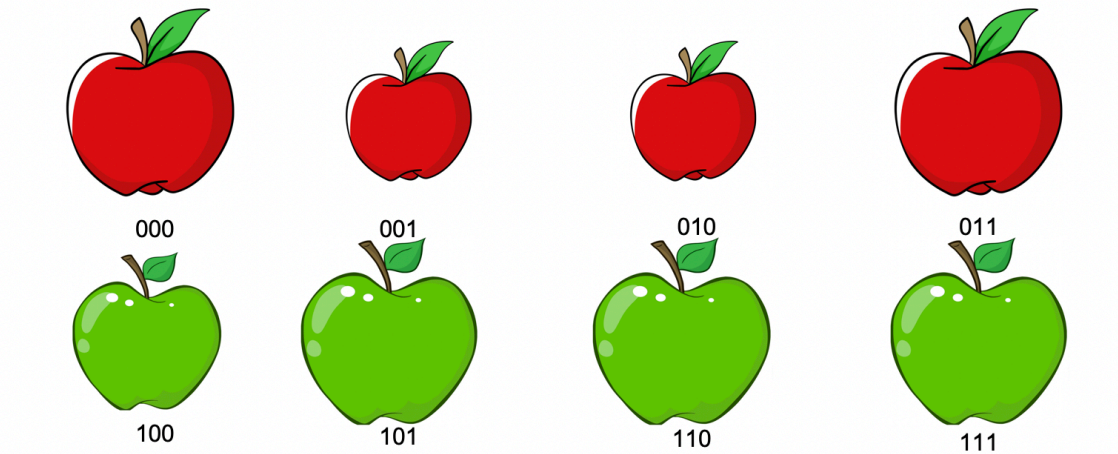

Secret:
Which apple is poisonous?

Before:
Each apple can be poisonous.
So we have 3 bit information.

Clue 1:
The poisonous apple is green.
Only four apples can be poisonous.
So we have 2 bit information.
Leaked = 3 bit – 2 bit = 1 bit

Eight Apples

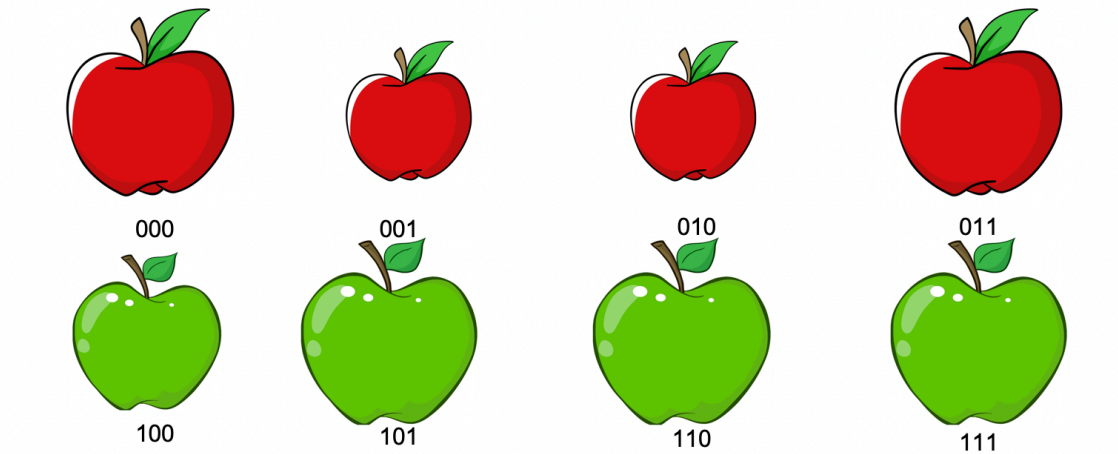

Secret:
Which apple is poisonous?

Before:
Each apple can be poisonous.
So we have 3 bit information.

Clue 1:
The poisonous apple is green.
Only four apples can be poisonous.
So we have 2 bit information.
Leaked = 3 bit – 2 bit = 1 bit

Eight Apples

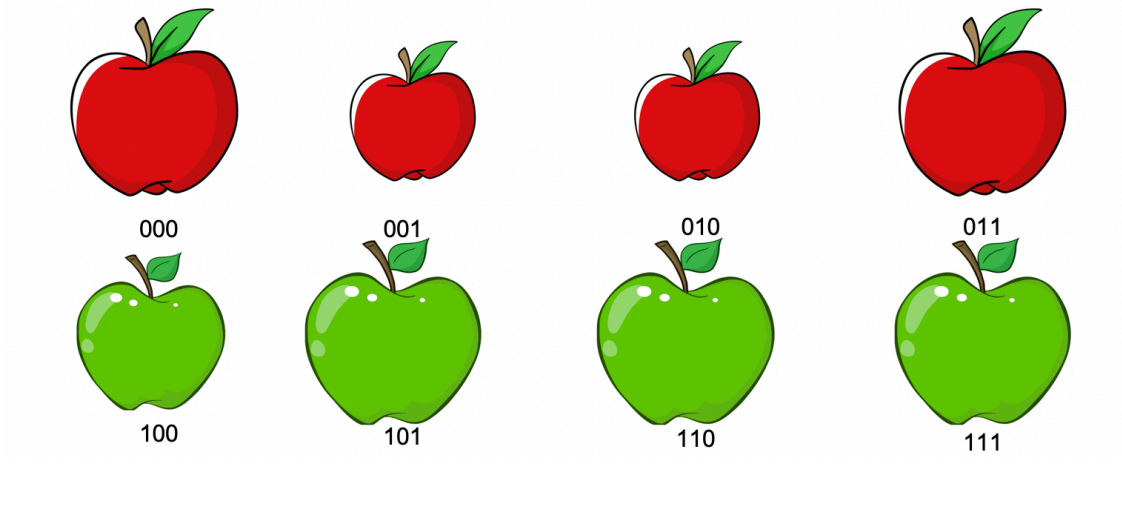

Number 100 apple is poisonous.

Secret:
Which apple is poisonous?

Before:
Each apple can be poisonous.
So we have 3 bit information.

Clue 1:
The poisonous apple is green.
Only four apples can be poisonous.
So we have 2 bit information.
Leaked = 3 bit – 2 bit = 1 bit

Clue 2:
The poisonous apple is small.

An example

```
unsigned char k1, k2;

if(k1 - k2 > 0)          // Leakage 1
    A();
if(k2 + k1 < 256)        // Leakage 2
    B();
```

An attacker:

1. knows the code runs A(), B()

Secret:
The value of k1 and k2
Before:
k1 and k2 can be any value from 0~255

Clue 1: k1 - k2 > 0
Clue 2: k1 + k2 < 256

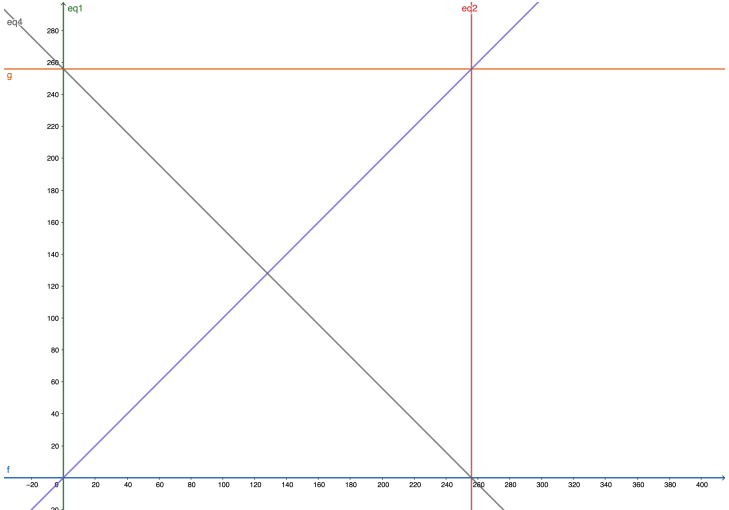

Information Leakage Definition

Definition. *Given a program P with the input set K, an adversary has the observation o when the input k ∈ K. We denote it as*

$$P(k) = 0$$

The leakage L_{ko} based on the observation is

$$L_{ko} = log_2|K| - log_2|K^o|$$

where

$$K_o = \{k' \mid k' \in K \text{ and } P(k') = o\}$$

Basic idea: counting the number of observations that have the same memory-access pattern.

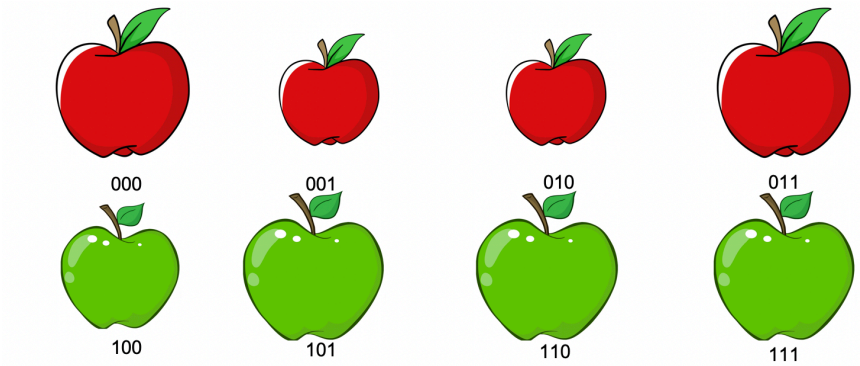

Clue 1:
The poisonous apple is green.

|K| = 8 $|K^o|$ = 4

L = log8 – log 4 = 1 bit

Information Leakage Definition

Definition. *Given a program P with the input set K, an adversary has the observation o when the input k ∈ K. We denote it as*

$$P(k) = 0$$

The leakage L_{ko} based on the observation is

$$L_{ko} = log_2|K| - log_2|K^o|$$

where

$$K_o = \{k' \mid k' \in K \text{ and } P(k') = o\}$$

Basic idea: counting the number of observations that have the same memory-access pattern.

```
// Dummy password checker
unsigned char key = input();
// key = [0 ... 255]
if(key = 128)
   A(); // branch 1
else if (key < 64)
   B(); // branch 2
else if (key < 128)
   C(); // branch 3
else
   D(); // branch 4
```

Example:

$|K| = 256 \quad |K_1| = 1 \quad |K_2| = 64$

$L_{ko1} = log_2 256 - log_2 1 = 8\ bits$

$L_{ko2} = log_2 256 - log_2 64 = 2\ bits$

Overview

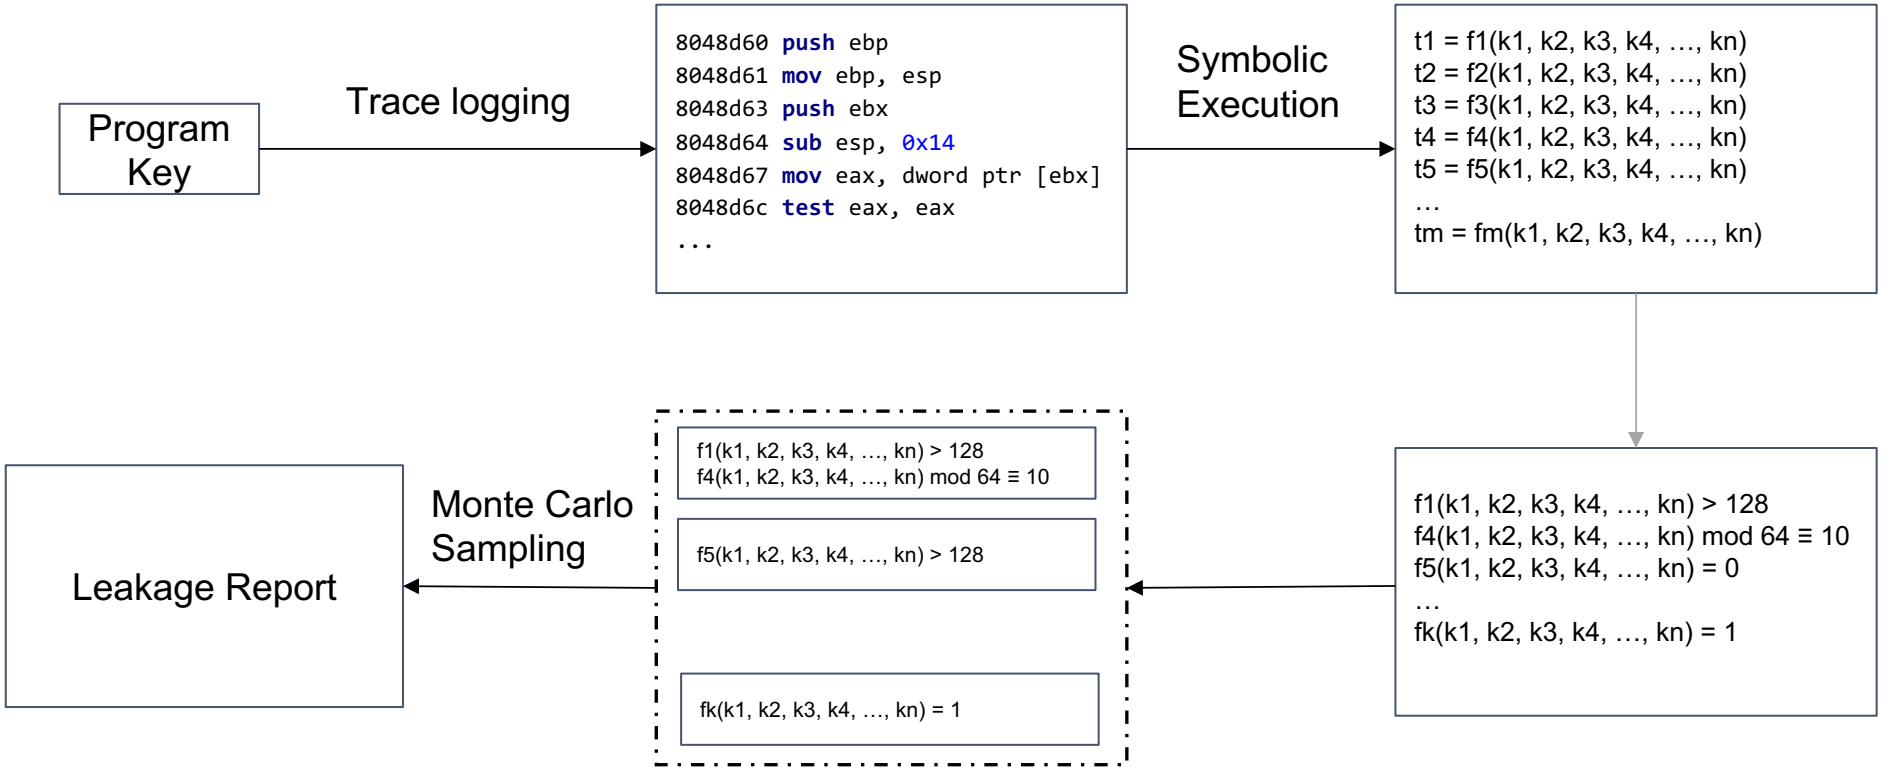

Figure 1. TBD architecture

Challenges

- Information Leakage Definition
  - Existing work: Shannon entropy and mutual information
  - Number of different observations
- Leakage Dependence
  - Real-world applications may have multiple information leakages sites
  - Some leakages are dependent
- Scalability
  - We want to quantify information leakages in real-world applications
  - e.g. RSA implementations OpenSSL have thousands of lines -> 1 million instructions
  - The performance is important

Challenge 1: Information Leakage Definition

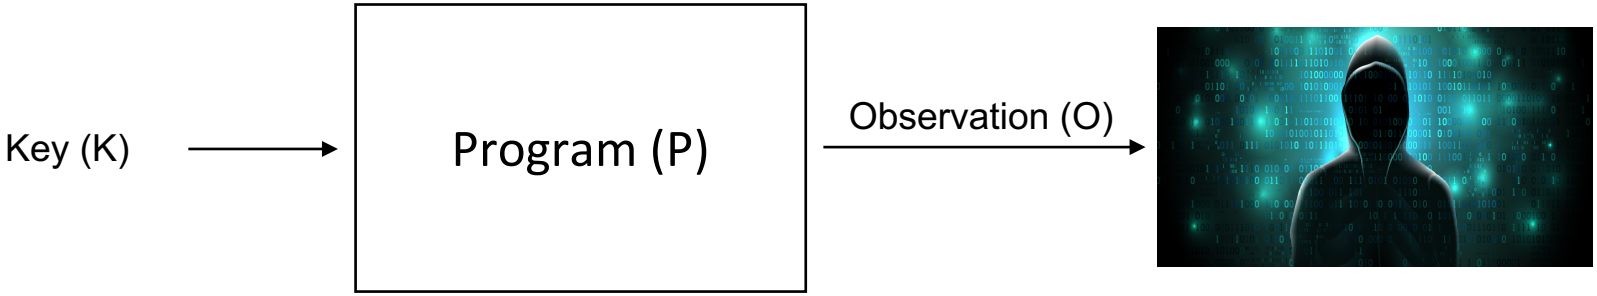

Figure 2. Observation Model

Observations:

1. Secret-dependent control-flow transfers
2. Secret-dependent memory access

Previous Information Leakage Definition

Mutual Information (MI):

$$\mathrm{MI} = \mathrm{I}(\mathrm{K};\mathrm{O}) = \mathrm{H}(\mathrm{O}) - \mathrm{H}(\mathrm{O}|\mathrm{K}) = \mathrm{H}(\mathrm{O}) = \sum_{\mathrm{o_i} \in \mathrm{O}} \mathrm{P}(\mathrm{o_i})\log_2 \mathrm{P}(\mathrm{o_i})$$

Maximal Leakage (ML):

$$\mathrm{ML} = \log_2 |\mathrm{O}|$$

```
unsigned char key = input();
// key = [0 ... 255]
if(key = 128)
   A(); // branch 1
else if (key < 64)
   B(); // branch 2
else if (key < 128)
   C(); // branch 3
else
   D(); // branch 4
```

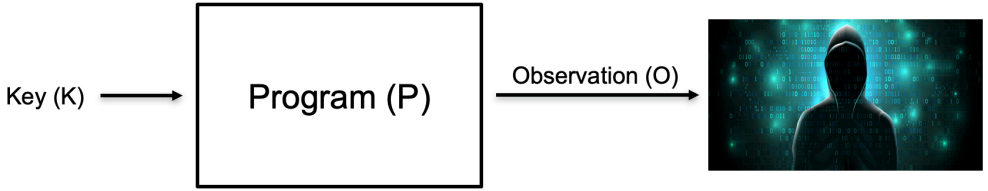

Branch	1	2	3	4
Possibility	$\frac{1}{256}$	$\frac{64}{256}$	$\frac{63}{256}$	$\frac{128}{256}$

$$MI = \frac{1}{256}log_2\frac{1}{256} + \frac{64}{256}log_2\frac{64}{256} + \frac{63}{256}log_2\frac{63}{256} + \frac{128}{256}log_2\frac{128}{256} = 1.7bits$$

$$ML = log_2 4 = 2bits$$

Challenge 1: Information Definition

```
// Dummy password checker
unsigned char key = input();
// key = [0 ... 255]
if(key = 128)
   A(); // branch 1
else if (key < 64)
   B(); // branch 2
else if (key < 128)
   C(); // branch 3
else
   D(); // branch 4
```

Suppose an attacker knows the code run branch 1, then he knows the key equals to 128.

Mutual Information: 1.7 bits
Max Leakage: 2 bits

Problem: Both methods neglect the input key and give an "average" estimate.

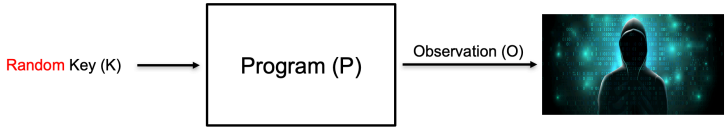

Figure 3. Mutual Information

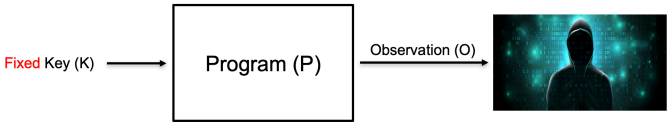

Figure 4. Real Attack

Challenge 2: Leakage Dependence

- Suppose one program has two leakage sites A and B.
  - Independent Leakage
  - Mutual Exclusive Leakage
  - Dependent Leakage

```
unsigned char k1, k2, t1, t2;
t1 = k1 + 2*k2;
t2 = 2*k1 - k2;
if(t2 + t1 > 127)  // Leakage A
    A();
if(t2 - t1 > 0)    // Leakage B
    B();
```

Leakage A:

$$\begin{cases} 0 \le k_1 \\ k_1 \le 255 \\ 3k_1 + k_2 > 127 \end{cases}$$

Leakage B:

$$\begin{cases} 0 \le k_2 \\ k_2 \le 255 \\ k_1 - 3k_2 > 0 \end{cases}$$

Leakage A and B :

$$\begin{cases} 0 \le k_1 \\ k_1 \le 255 \\ 3k_1 + k_2 > 127 \\ k_1 - 3k_2 > 0 \end{cases}$$

if $L_{(k1k2)A} + L_{(k1+k2)B} = L_{(k1+k2)AB}$,
then A and B are independent

Challenge 3: Scalability

We want to quantify information leakages from real-world applications.

The performance suffers from the following aspects:

- Symbolic Execution
  - IR explosion
- Monte Carlo Sampling
  - #P problem

Challenge 3: Scalability (Symbolic Execution)

- Why symbolic execution (SE) is slow?
  - Path Explosion ✔
  - Constraint Solving ✔
  - Intermediate Languages (IR) ?

Why symbolic execution uses IR?

- Support many architecture platform.
  - X86, ARM, MIPS
- Easy to implement
  - X86, more than 1000 instructions
  - Side effects

test eax, eax →

```
STR  R_EAX:32,     , V_00:32
STR  0:1,          , R_CF:1
AND  V_00:32,  ff:8, V_01:8
SHR  V_01:8,    7:8, V_02:8
SHR  V_01:8,    6:8, V_03:8
XOR  V_02:8, V_03:8, V_04:8
SHR  V_01:8,    5:8, V_05:8
SHR  V_01:8,    4:8, V_06:8
XOR  V_05:8, V_06:8, V_07:8
XOR  V_04:8, V_07:8, V_08:8
SHR  V_01:8,    3:8, V_09:8
SHR  V_01:8,    2:8, V_10:8
XOR  V_09:8, V_10:8, V_11:8
SHR  V_01:8,    1:8, V_12:8
XOR  V_12:8, V_01:8, V_13:8
XOR  V_11:8, V_13:8, V_14:8
XOR  V_08:8, V_14:8, V_15:8
AND  V_15:8,    1:1, V_16:1
NOT  V_16:1,       , R_PF:1
STR     0:1,       , R_AF:1
EQ  V_00:32,   0:32, R_ZF:1
SHR V_00:32,  1f:32, V_17:32
AND    1:32,V_17:32, V_18:32
EQ     1:32,V_18:32, R_SF:1
STR     0:1,       , R_OF:1
```

Solution: instruction-level symbolic execution

QSYM: A Practical Concolic Execution Engine Tailored for Hybrid Fuzzing. Insu Yun, Sangho Lee, Meng Xu, Yeongjin Jang, and Taesoo Kim.In Proceedings of the 27th USENIX Security Symposium (Security 2018),Baltimore, MD, August 2018.

Challenge 3: Scalability (Monte Carlo Sampling)

- The problem is #P-Hard.
- The number of satisfying assignments might be exponentially small.
  - e.g., $F(K) = \begin{cases} k_1 = 1 \\ 120 < k_2 < 123 \\ k_3 = 3 \\ k_4 = 4 \\ k_5 = 5 \end{cases} \quad k_1, k_2, k_3, k_4, k_5 \in [0, 255]$

Total search space: 2^{40}
Only two satisfying assignments

Solution: Split the constraint into several independent small groups

Implementation

- 16k LoC of C++11
- Trace Collection
  - Intel Pin Tool
- Symbolic Execution
  - arithmetic, bitwise, logical, control-transfer
  - Not support: AVX, floating number, AES-NI
- Information Leakage Computation
- Data flow analysis
- Debug Information Parser

Evaluation (Not finished)

- Real-world cryptosystems:
  - OpenSSL 0.9.7 1.1.1
  - mbedTLS 2.5 2.15
  - Libjpeg
- Performance
  - 8s found 8 secret-dependent memory access for the DES in mbedTLS 2.5

Demo

Summary

Identify and quantify address-based side-channels leakages precisely

- A trace-based method that models each leakage site with math formulas.
- We quantify the information leakages based on the number of secrets that satisfy the constrain.
- Most of leakages found by recent works are hard to be exploited. (Need to be confirmed)